

CBOR (RFC 7049)

Concise Binary Object Representation

IETF94 CBOR lightning tutorial

Carsten Bormann, 2015-11-01

CBOR: Agenda


- What is it, and when might I want it?
- How does it work?
- How do I work with it?

CBOR: Agenda

- What is it, and when might I want it?
- How does it work?
- How do I work with it?

History of Data Formats

- **Ad Hoc**
- **Database Model**
- **Document Model**
- **Programming Language Model**




```

S:<?xml version="1.0" encoding="UTF-8" standalone="no"?>
S:<epp xmlns="urn:ietf:params:xml:ns:epp-1.0"
S:  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
S:  xsi:schemaLocation="urn:ietf:params:xml:ns:epp-1.0
S:    epp-1.0.xsd">
S:<response>
S:  <result code="1000">
S:    <msg>Command completed successfully</msg>
S:  </result>
S:  <resData>
S:    <domain:infData
S:      xmlns:domain="urn:ietf:params:xml:ns:domain-1.0"
S:      xsi:schemaLocation="urn:ietf:params:xml:ns:domain-1.0
S:        domain-1.0.xsd">
S:        <domain:name>3.8.0.0.6.9.2.3.6.1.4.4.e164.arpa</domain:name>
S:        <domain:roid>EXAMPLE1-REP</domain:roid>
S:        <domain:status s="ok"/>
S:        <domain:registrant>jd1234</domain:registrant>
S:        <domain:contact type="admin">sh8013</domain:contact>
S:        <domain:contact type="tech">sh8013</domain:contact>
S:        <domain:ns>
S:          <domain:hostObj>ns1.example.com</domain:hostObj>
S:          <domain:hostObj>ns2.example.com</domain:hostObj>
S:        </domain:ns>
S:        <domain:hostObj>example.com</domain:hostObj>
S:        <domain:hostObj>ns2.example.com</domain:hostObj>
S:        <domain:clID>ClientX</domain:clID>
S:        <domain:crID>ClientY</domain:crID>
S:        <domain:crDate>1999-04-03T22:00:00.0Z</domain:crDate>
S:        <domain:upID>ClientX</domain:upID>
S:        <domain:upDate>1999-12-03T09:00:00.0Z</domain:upDate>
S:        <domain:exDate>2005-04-03T22:00:00.0Z</domain:exDate>
S:        <domain:trDate>2000-04-08T09:00:00.0Z</domain:trDate>
S:        <domain:authInfo>
S:          <domain:pw>2fooBAR</domain:pw>
S:        </domain:authInfo>
S:      </domain:infData>
S:    </resData>
S:    <extension>
S:      <e164:infData xmlns:e164="urn:ietf:params:xml:ns:e164epp-1.0"
S:        xsi:schemaLocation="urn:ietf:params:xml:ns:e164epp-1.0
S:          e164epp-1.0.xsd">
S:        <e164:naptr>
S:          <e164:order>10</e164:order>
S:          <e164:pref>100</e164:pref>
S:          <e164:flags>u</e164:flags>

```

XML

```

type="idmef:file-permission"
use="required" />
</xsd:complexType>

<xsd:complexType name="FileAccess">
  <xsd:sequence>
    <xsd:element name="UserId"
      type="idmef:UserId" />
    <xsd:element name="permission"
      type="idmef:Permission"
      minOccurs="1"
      maxOccurs="unbounded" />
  </xsd:sequence>
</xsd:complexType>

<xsd:complexType name="Inode">
  <xsd:sequence>
    <xsd:element name="change-time"
      type="xsd:string"
      minOccurs="0"
      maxOccurs="1" />
    <xsd:element name="number"
      type="xsd:string" />
    <xsd:element name="major-device"
      type="xsd:string" />
    <xsd:element name="minor-device"
      type="xsd:string" />
  </xsd:sequence>
  <xsd:sequence minOccurs="0" maxOccurs="1">
    <xsd:element name="c-major-device"
      type="xsd:string" />
    <xsd:element name="c-minor-device"
      type="xsd:string" />
  </xsd:sequence>
</xsd:complexType>

<xsd:complexType name="Linkage">
  <xsd:choice>
    <xsd:sequence>
      <xsd:element name="name" type="xsd:string" />
      <xsd:element name="path" type="xsd:string" />
    </xsd:sequence>
    <xsd:element name="File" type="idmef:File" />
  </xsd:choice>
  <xsd:attribute name="category"
    type="idmef:linkage-category"

```

XSD

JSON

- **JavaScript Object Notation**
- **Minimal**
- **Textual**
- **Subset of JavaScript**

Values

- **Strings**
- **Numbers**
- **Booleans**

- **Objects**
- **Arrays**

- `null`

Array

```
["Sunday", "Monday",  
 "Tuesday", "Wednesday",  
 "Thursday", "Friday",  
 "Saturday"]
```

```
[  
    [0, -1, 0],  
    [1, 0, 0],  
    [0, 0, 1]  
]
```

Object

```
{  
  "name":      "Jack B. Nimble",  
  "at large":  true,  
  "grade":     "A",  
  "format": {  
    "type":     "rect",  
    "width":    1920,  
    "height":   1080,  
    "interlace": false,  
    "framerate": 24  
  }  
}
```

~~Object~~ Map

```
{  
  "name": "Jack B. Nimble",  
  "at large": true,  
  "grade": "A",  
  "format": {  
    "type": "rect",  
    "width": 1920,  
    "height": 1080,  
    "interlace": false,  
    "framerate": 24  
  }  
}
```

Application usage of JSON

- **No schema needed** for parsing
- Representation types:
bool, number, string, null; **map**, **array**
- Semantics:
 - **struct**: key names in map (vs. **table** use of maps)
 - **record**: position in array (vs. **vector** use of arrays)
- Extensibility: new keys in a key/value struct
 - Ignore what you don't understand

JSON limitations

- No **binary** data (byte strings)
- Numbers are in **decimal**, some parsing required
- Format requires copying:
 - **Escaping** for strings
 - Base64 for binary
- **No extensibility** (e.g., date format?)
- Interoperability **issues**
 - I-JSON further reduces functionality (RFC 7493)

	Character-based	Concise Binary
Document-Oriented	XML	EXI
Data-Oriented	JSON	???

BSON and friends

- Lots of “binary JSON” proposals
- Often optimized for data at rest, not protocol use (BSON → MongoDB)
- Most are **more** complex than JSON

Why a new binary object format?

- Different design goals from current formats
 - stated up front in the document
- Extremely **small code size**
 - for work on constrained node networks
- Reasonably **compact data size**
 - but no compression or even bit-fiddling
- Useful to any protocol or application that *likes* the design goals

Concise Binary Object Representation (CBOR)



“Sea Boar”

	Character-based	Concise Binary
Document-Oriented	XML	EXI
Data-Oriented	JSON	CBOR

Design goals (1 of 2)

1. unambiguously encode most **common data formats** (such as JSON-like data) used in Internet standards
2. **compact implementation** possible for encoder and decoder
3. able to parse **without a schema description.**

Design goals (2 of 2)

4. Serialization reasonably **compact**, but data compactness **secondary** to implementation compactness
5. applicable to both **constrained nodes** and **high-volume applications**
6. support all **JSON** data types, conversion to and from JSON
7. **extensible**, with the extended data being able to be parsed by earlier parsers

2013-09-13: CBOR RFC

- “Concise Binary Object Representation”:
JSON equivalent for constrained nodes
 - start from JSON data model (no schema needed)
 - add binary data, extensibility (“tags”)
 - concise binary encoding (byte-oriented, counting objects)
 - add diagnostic notation
- Done without a WG (with APPSAWG support)

<http://cbor.io>

CBOR

RFC 7049 Concise Binary Object Representation

“The Concise Binary Object Representation (CBOR) is a data format whose design goals include the possibility of extremely small code size, fairly small message size, and extensibility without the need for version negotiation.”

JSON data model

CBOR is based on the wildly successful JSON data model: numbers, strings, arrays, maps (called objects in JSON), and a few values such as false, true, and null.

No Schema needed

Embracing binary

Some applications that would like to use JSON need to transport binary data, such as encryption keys, graphic data, or sensor values. In JSON, these data need to be encoded (usually in base64 format), adding complexity and bulk.

Concise encoding

Stable format

CBOR is defined in an Internet Standards Document, [RFC 7049](#). The format has been designed to be stable for decades.

Extensible

To be able grow with its applications and to

CBOR: Agenda

- What is it, and when might I want it?
- How does it work?
- How do I work with it?

CBOR vs. “binary JSONs”

- Encoding [1, [2, 3]]: compact | stream

ASN.1 BER*	30 0b 02 01 01 30 06 02 01 02 02 01 03	30 80 02 01 01 30 06 02 01 02 02 01 03 00 00
MessagePack	92 01 92 02 03	
BSON	22 00 00 00 10 30 00 01 00 00 00 04 31 00 13 00 00 00 10 30 00 02 00 00 00 10 31 00 03 00 00 00 00 00	
UBJSON	61 02 42 01 61 02 42 02 42 03	61 ff 42 01 61 02 42 02 42 03 45*
CBOR	82 01 82 02 03	9f 01 82 02 03 ff

Very quick overview of the format

- Initial byte: **major type** (3 bits) and **additional information** (5 bits: immediate value or length information)
- Eight major types:
 - unsigned (0) and negative (1) **integers**
 - **byte** strings (2), **UTF-8** strings (3)
 - **arrays** (4), **maps** (5)
 - optional **tagging** (6) and **simple types** (7) (floating point, Booleans, etc.)

Additional information

- 5 bits
 - 0..23: immediate value
 - 24..27: 1, 2, 4, 8 bytes value follow
 - 28..30: *reserved*
 - 31: indefinite length
 - terminated only by 0xFF in place of data item
- Generates unsigned integer:
 - **Value** for mt 0, 1 (unsigned/neg integers), 7 (“simple”)
 - **Length** (in bytes) for mt 2, 3 (byte/text strings)
 - **Count** (in items) for mt 4, 5 (array, map)
 - **Tag** value for mt 6

Major types 6 and 7

- mt 7:
 - **special** values for $a_i = 0..24$
 - false, true, null, undef
 - IANA registry for more
 - $a_i = 25, 26, 27$: **IEEE floats**
 - in 16 (“half”), 32 (“single”), and 64 (“double”) bits
- mt 6: semantic **tagging** for things like dates, arbitrary-length bignums, and decimal fractions

Tags

- A Tag contains **one** data item
- 0: RFC 3339 (~ ISO 8601) **text string** date/time
- 1: UNIX time (**number** relative to 1970-01-01)
- 2/3: bignum (**byte string** encodes unsigned)
- 4: [**exp, mant**] (decimal fraction)
- 5: [**exp, mant**] (binary fraction, “bigfloat”)
- 21..23: expected conversion of **byte string**
- 24: nested CBOR data item in **byte string**
- 32...: URI, base64[url], regexp, mime (**text strings**)

New Tags

- Anyone can register a tag (IANA)
 - 0..23: Standards action
 - 24..255: Specification required
 - 256..18446744073709551615: FCFS
- 25/256: stringref for simple compression
- 28/29: value sharing (beyond trees)
- 26/27: constructed object (Perl/generic)
- 22098: Perl reference (“indirection”)

Examples

- Lots of examples in RFC (making use of JSON-like “diagnostic notation”)
- $0 \rightarrow 0x00$, $1 \rightarrow 0x01$, $23 \rightarrow 0x17$, $24 \rightarrow 0x1818$
- $100 \rightarrow 0x1864$, $1000 \rightarrow 0x1903e8$, $1000000 \rightarrow 0x1a000f4240$
- $18446744073709551615 \rightarrow 0x1b\text{ffffffffffffffff}$, $18446744073709551616 \rightarrow 0xc249010000000000000000$
- $-1 \rightarrow 0x20$, $-10 \rightarrow 0x29$, $-100 \rightarrow 0x3863$, $-1000 \rightarrow 0x3903e7$
- $1.0 \rightarrow 0xf93c00$, $1.1 \rightarrow 0xfb3ff1999999999999a$, $1.5 \rightarrow 0xf93e00$
- $\text{Infinity} \rightarrow 0xf97c00$, $\text{NaN} \rightarrow 0xf97e00$, $-\text{Infinity} \rightarrow 0xf9fc00$
- $\text{false} \rightarrow 0xf4$, $\text{true} \rightarrow 0xf5$, $\text{null} \rightarrow 0xf6$
- $h'' \rightarrow 0x40$, $h'01020304' \rightarrow 0x4401020304$
- $"" \rightarrow 0x60$, $"a" \rightarrow 0x6161$, $"IETF" \rightarrow 0x6449455446$
- $[] \rightarrow 0x80$, $[1, 2, 3] \rightarrow 0x83010203$, $[1, [2, 3], [4, 5]] \rightarrow 0x8301820203820405$
- $\{\} \rightarrow 0xa0$, $\{1: 2, 3: 4\} \rightarrow 0xa201020304$, $\{"a": 1, "b": [2, 3]\} \rightarrow 0xa26161016162820203$

CBOR: Agenda

- What is it, and when might I want it?
- How does it work?
- How do I work with it?

<http://cbor.me>: CBOR playground

- Convert back and forth between **diagnostic notation** (~JSON) and binary encoding

CBOR

Diagnostic 

 **5** Bytes

[1, [2, 3]]

82	# array(2)
01	# unsigned(1)
82	# array(2)
02	# unsigned(2)
03	# unsigned(3)

Implementations

- Parsing/generating CBOR easier than interfacing with application
- Minimal implementation: 822 bytes of ARM code
- Different integration models, different languages
- > 25 implementations (after first two years)

The screenshot displays the 'Implementations' section of the <http://cbor.io> website. It is organized into three columns, each with a header for a group of languages. Each language entry includes a brief description, installation instructions, and a 'View details' link.

JavaScript	Lua	C#, Java
JavaScript implementations are becoming available both for in-browser use and for node.js. Browser A CBOR object can be installed via <code>bower install cbor</code> and used as an AMD module or global object in the browser e.g. in combination with Websockets... View details » node.js ... and the server side for that might be written using node.js: install via: <code>npm install cbor</code> View details »	Lua-cbor is a pure Lua implementation of CBOR for Lua 5.1–5.3, which utilizes struct packing and bitwise operations if available: View details »	A rather comprehensive implementation that addresses arbitrary precision arithmetic is available in both a C# and a Java version. View details »
PHP API: <code>\CBOR\CBOREncoder::encode(\$target)</code> and <code>\CBOR\CBORDecoder::decode(\$encoded_d)</code> View details »	Python Install a high-speed implementation via pypi: <code>pip install cbor</code> View details » Flynn's' simple API is inspired by existing Python serialisation modules like json and pickle: View details »	Java A Java implementation as part of the popular Jackson JSON library is at: View Details » A Java 7 implementation focusing on test coverage and a clean separation of model, encoder and decoder is at: View Details » JACOB, a small CBOR encoder and decoder implemented in plain Java is at: View Details »
Go An early Go implementation that feels like the JSON library: View details » Another, more full-grown Go implementation: View details » Most recently, a comprehensive, high-performance implementation has become available as part of a larger set of data representation format en- and decoders: View details »	Perl Install a comprehensive implementation tailored to Perl's many features via: <code>cpan CBOR::XS</code> You'll like the performance data... View details »	C, C++ A CBOR implementation in C is part of the RIOT operating system for constrained nodes: View Details » A C implementation for highly constrained nodes, which achieves a full CBOR decoder in 880 bytes of ARM code (and now also includes an encoder), has recently become available. View Details » A basic C++ implementation is also available: View Details »
Rust A Rust implementation is available that works with Cargo and is on crates.io : View details » Another Rust implementation has also become available recently on crates.io : View details »	Ruby A high-speed implementation has been derived from the MessagePack implementation for Ruby. Installation: <code>gem install cbor</code> View details » Ruby bindings for <code>libcbor</code> are now available. Installation: <code>gem install libcbor</code> View details »	Erlang, Elixir cbor-erlang is a recent implementation in Erlang: View details » An older Elixir implementation is also available: <code>expm spec excbor --format scm sh</code> Or look at the source: View details »
	Haskell Now on hackage : View details »	D A compact D implementation with a Dub package : View Details »

<http://cbor.io>

CBOR and CDDL

- CBOR takeup within IETF is increasing.
How to write specs?
- CDDL: CBOR **Data Definition** Language
<https://tools.ietf.org/html/draft-greevenbosch-appsawg-cbor-cddl-07>
 - The best of ABNF, Relax-NG, JSON Content Rules
 - Rough tool available: **gem install cddl**
 - **Generate** example instances (CBOR or JSON)
 - **Check** instances against the definition

How **RFC 7071** would have looked like in CDDL

reputation-object = {	; This is a map (JSON object)
application: text	; text string (vs. binary)
reputons: [* reputon]	; Array of 0- ∞ reputons
}	

reputon = {	; Another map (JSON object)
rater: text	
assertion: text	
rated: text	
rating: float16	; OK, float16 is a CBORism
? confidence: float16	; optional...
? normal-rating: float16	
? sample-size: uint	; unsigned integer
? generated: uint	
? expires: uint	
* text => any	; 0- ∞ , express extensibility
}	

GRASP

- Generic Autonomic Signaling Protocol (GRASP)
- For once, try not to invent another TLV format: just use CBOR
- Messages are arrays, with type, id, option:
 `message /= [MESSAGE_TYPE, session-id, *option]`
 `MESSAGE_TYPE = 123 ; a defined constant`
 `session-id = 0..16777215`
 `; option is one of the options defined below`
- Options are arrays, again:
 `option /= waiting-time-option`
 `waiting-time-option =`
 `[0_WAITING, waiting-time]`
 `0_WAITING = 456 ; a defined constant`
 `waiting-time = 0..4294967295 ; in milliseconds`

Where from here?

- RFC 7049
- <http://cbor.io>
- cbor@ietf.org
- <http://tools.ietf.org/html/cddl>